

DataEng S24: Data Storage In-class Assignment

This week you'll gain experience with various ways to load data into a database.

Make a copy of this document and use it to record your results. Store a PDF copy of the document in your git repository along with any needed code before submitting for this week.

The data set for this week is US Census data from 2015. The United States conducts a full census of every household every 10 years (last one was in 2020), but much of the detailed census data comes during the intervening years when the Census Bureau conducts its detailed American Community Survey (ACS) of a randomly selected sample of approximately 3.5 million households each year. The resulting data gives a more detailed view of many factors of American life and the composition of households.

[ACS Census Tract Data for 2015 \(part 1\)](#)

[ACS Census Tract Data for 2015 \(part 2\)](#)

Your job is to load the 2015 data set (approximately 74000 rows divided into two parts). You'll configure a postgres DBMS on a new GCP virtual machine, and then load the two parts of the data using multiple loading methods, comparing the cost of each method. Load the two parts of the data separately so that you gain experience with the issues related to incremental loading of data.

Note that the goal here is not to achieve the fastest load times. Instead, your goal should be to observe and understand a variety of data loading methods. If you start to run out of time, then skip to part I (copy_from) to be sure to get experience with what is probably the fastest way to load bulk data into a Postgres server.

Please highlight your responses so that we can more easily see them

A. [MUST] Discussion Question (discuss as a group near the beginning of the week. Note your own response in this space):

Do you have any experience with ingesting bulk data into a DBMS? If yes, then describe your experience, especially what method was used to input the data into the database. If no, then describe how you might ingest daily incremental breadcrumb data for your class project.

Response:

Yes, I have experience with bulk data ingestion into a DBMS. In a previous project, we imported a large dataset into a relational database using CSV uploads and bulk insert operations. After organizing the data into clean CSV files, we used PostgreSQL's COPY command for efficient loading into the database tables.

For ingesting daily incremental breadcrumb data in a class project:

1. Extract the daily data from the source.
2. Transform the data to match the database schema, including cleaning, filtering, and ensuring consistency.
3. Load the transformed data incrementally into the database.
4. Monitor and validate the ingestion process continuously to ensure accuracy and completeness.

Submit: submit your assignment by this Friday at 10pm

B. [MUST] Configure the Database

1. Create a new GCP virtual machine for this week's work (medium size or larger).
2. Follow the steps listed in the [Installing and Configuring a PostgreSQL server](#) instructions. To keep your project work clean, use a new, different vm and delete the vm when finished with this assignment.
3. Also running the following commands on your VM will help to configure the python module "psycopg2" which you will use to connect to your postgres database:

4.

```
sudo apt install python3 python3-dev python3-venv
sudo apt-get install python3-pip
pip3 install psycopg2-binary
```

```
Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
Last login: Fri May 10 21:21:16 2024 from 35.235.241.65
dileepkb@instance-20240510-203806:~$ ls
AL2015_1.csv  OR2015_2.csv  acs2015_census_tract_data_part1.csv  acs2015_census_tract_data_part2.csv  myvenv  venv
dileepkb@instance-20240510-203806:~$ cd myvenv
dileepkb@instance-20240510-203806:~/myvenv$ source venv/bin/activate
(venv) dileepkb@instance-20240510-203806:~/myvenv$ cd ..
(venv) dileepkb@instance-20240510-203806:~$ ls
AL2015_1.csv  OR2015_2.csv  acs2015_census_tract_data_part1.csv  acs2015_census_tract_data_part2.csv  myvenv  venv
(venv) dileepkb@instance-20240510-203806:~$ sudo vim load_inserts.py
```

C. [MUST] Prepare Data for Loading

1. Copy/upload both ACS Census Tract data files to your VM
2. Create a small test sample by running a linux command similar to the following. The small test sample will help you to quickly test any code that you write before running your code with the full dataset.

```
head -1 acs2015_census_tract_data_part1.csv > AL2015_1.csv
```

```
head -1 acs2015_census_tract_data_part2.csv > OR2015_2.csv
grep Alabama acs2015_census_tract_data_part1.csv >> AL2015_1.csv
grep Oregon acs2015_census_tract_data_part2.csv >> OR2015_2.csv
```

The first two commands copy the headers to the sample files and the next command appends Alabama and Oregon 2015 data to the sample files. This should produce files with fewer than 1500 records per file. Use these sample files to save time during testing.

```
dileepkb@instance-20240510-203806:~$ ls
AL2015_1.csv  acs2015_census_tract_data_part1.csv  load_inserts.py  venv
OR2015_2.csv  acs2015_census_tract_data_part2.csv  myvenv
```

3. Write a python program that connects to your postgres database and creates your main census data table. Start with this example code: [load_inserts.py](#). For example:

```
python3 load_inserts.py -d ./AL2015_1.csv -c
python3 load_inserts.py -d ./OR2015_2.csv

(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part1.csv
Created CensusData
Loading 36844 rows
Finished Loading. Elapsed Time: 37.48 seconds

dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv
readdata: reading from File: ./acs2015_census_tract_data_part2.csv
Loading 37157 rows
Finished Loading. Elapsed Time: 38.57 seconds
```

D. [MUST] Baseline - Simple INSERT

The tried and true SQL command [INSERT INTO ...](#) is the most basic way to insert data into a SQL database, and often it is the best choice for small amounts of data, production databases and other situations in which you need to maintain performance and ACID properties of the updated table.

The `load_inserts.py` program shows how to use simple INSERTs to load data into a database. It is possibly the slowest way to load large amounts of data. For me, it takes approximately 1 second for the Oregon sample and nearly 60 seconds to load each part of the acs data.

Take the program and try it with both of the test samples and both parts of the ACS data set. Fill in the appropriate information in the table below (see part J).

After loading each part of the ACS data, try out a few validations to check that the data is loaded correctly. Use the following SQL queries and/or create more of your own:

- After loading part1,
 - The number of states in the database should be 24
 - `select count(distinct state) from censusdata ;`
 - The state of Portland is not found in the database

- select 1 from censusdata where state = 'Portland' limit 1 ;
- The state of Oregon is not found in the database
 - select count(*) from censusdata where state = 'Oregon' ;
- The state of Iowa is found in the database
 - select 1 from censusdata where state = 'Iowa' limit 1 ;
- There are 99 counties in Iowa
 - select count(distinct county) from censusdata where state = 'Iowa' ;

```
psql (15.6 (Debian 15.6-0+deb12u1))
Type "help" for help.

postgres=# select count(distinct state) from censusdata ;
 count
-----
      24
(1 row)

postgres=# select 1 from censusdata where state = 'Portland' limit 1 ;
?column?
-----
(0 rows)

postgres=# select count(*) from censusdata where state = 'Oregon' ;
 count
-----
      0
(1 row)

postgres=# select 1 from censusdata where state = 'Iowa' limit 1 ;
?column?
-----
      1
(1 row)

postgres=# select count(distinct county) from censusdata where state = 'Iowa' ;
 count
-----
     99
(1 row)
```

- After loading part2,
 - The number of states in the database should be 52
 - select count(distinct state) from censusdata ;
 - The state of Portland is not found in the database
 - select 1 from censusdata where state = 'Portland' limit 1 ;
 - The state of Oregon is found in the database
 - select count(*) from censusdata where state = 'Oregon' ;

- There are 36 counties in Oregon
 - `select count(distinct county) from censusdata where state = 'Oregon' ;`
- The state of Iowa is found in the database
 - `select 1 from censusdata where state = 'Iowa' limit 1 ;`
- There are 99 counties in Iowa
 - `select count(distinct county) from censusdata where state = 'Iowa' ;`

```
postgres=# select count(distinct state) from censusdata ;
count
-----
      52
(1 row)

postgres=# select 1 from censusdata where state = 'Portland' limit 1 ;
?column?
-----
(0 rows)

postgres=# select count(*) from censusdata where state = 'Oregon' ;
count
-----
   834
(1 row)

postgres=# SELECT COUNT(DISTINCT county) FROM censusdata WHERE state = 'Oregon';
count
-----
    36
(1 row)

postgres=# select 1 from censusdata where state = 'Iowa' limit 1 ;
?column?
-----
      1
(1 row)

postgres=# select count(distinct county) from censusdata where state = 'Iowa' ;
count
-----
    99
(1 row)
```

E. [MUST] Disabling Indexes and Constraints

You might notice that the CensusData table has a Primary Key constraint and an additional index on the state name column. Indexes and constraints are helpful for query performance but these features can slow down load performance.

Modify your `load_inserts.py` program to delay creation of constraints/indexes until after the data set is loaded. Enter the resulting load time into the results table below. How much does this technique improve load performance?

```
(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part1.csv
Created CensusData
Loading 36844 rows
Finished Loading. Elapsed Time: 34.83 seconds
Added Primary Key and Index
(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part2.csv
Created CensusData
Loading 37157 rows
Finished Loading. Elapsed Time: 35.5 seconds
Added Primary Key and Index
```

F. [SHOULD] Disabling Autocommit

You might have noticed that the `load_inserts.py` program sets `autocommit=True` on the database connection. This makes loaded data available to DB queries immediately after each insert. But it also triggers transaction-related overhead operations (e.g., write ahead logging). It also allows readers of the database to view an incomplete set of data during the load.

Modify your `load_inserts.py` program to avoid setting `autocommit=True`

Enter the resulting load time into the results table below.

```
(venv) dileepkb@instance-20240510-203806:~$ nano load_inserts.py
(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part1.csv
Created CensusData
Loading 36844 rows
Finished Loading. Elapsed Time: 6.202 seconds
Added Primary Key and Index
(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part2.csv
Created CensusData
Loading 37157 rows
Finished Loading. Elapsed Time: 5.054 seconds
Added Primary Key and Index
```

G. [SHOULD] UNLOGGED table

By default, RDBMS tables incur overheads of write-ahead logging (WAL) such that the database outputs extra metadata about each row insert to a log file known as the Transaction Recovery Log (sometimes just called the WAL or “Write-Ahead Log”). The RDBMS uses that WAL data to recover the contents of the table if/when the RDBMS crashes. Crash Recovery is a great feature but it can slow down load performance.

You can avoid this extra WAL overhead by [loading to an UNLOGGED table](#). Modify your `load_inserts.py` program to load data to an UNLOGGED table. Then enhance `load_inserts.py` to use a SQL query to append the staging data to the main `CensusData` table. Then create the needed index and constraint.

```
(venv) dileepkb@instance-20240510-203806:~$ nano load_inserts.py
(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part1.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part1.csv
Created CensusData_staging
Loading 36844 rows into staging table
Finished Loading into staging table. Elapsed Time: 7.967 seconds
Appended staging data to main table
Added Primary Key and Index
(venv) dileepkb@instance-20240510-203806:~$ python3 load_inserts.py -d ./acs2015_census_tract_data_part2.csv -c
readdata: reading from File: ./acs2015_census_tract_data_part2.csv
Created CensusData_staging
Loading 37157 rows into staging table
Finished Loading into staging table. Elapsed Time: 7.95 seconds
```

H. [ASPIRE] Temp Tables and Memory Tuning

Next compare the above approach with loading the data to [a temporary table](#) (and copying from the temporary table to the CensusData table). Which approach works best for you?

The amount of memory used for temporary tables is default configured to only 8MB. Your VM has enough memory to allocate much more memory to temporary tables. Try allocating 256 MB (or more) to temporary tables. So update the [temp_buffers parameter](#) to allow the database to use more memory for your temporary table. Rerun your load experiments. Did it make a difference?

I. [MUST] Built In Facility (copy_from)

The number one rule of bulk loading is to pay attention to the native facilities provided by the DBMS system implementers. DBMS vendors often put great effort into providing purpose-built loading mechanisms that achieve high performance and scalability.

With a simple, one-server Postgres database, that facility is known as COPY, \copy, or for python programmers [copy_from](#). See the last section of [this blog post](#) for an example.

```
(venv) dileepkb@instance-20240510-203806:~$ python copy_part1.py -f /home/dileepkb/acs2015_census_tract_data_part1.csv
Table censusdata truncated.
Data loaded successfully. Elapsed Time: 0.3292 seconds
(venv) dileepkb@instance-20240510-203806:~$ nano copy_part2.py
(venv) dileepkb@instance-20240510-203806:~$ python copy_part2.py -f /home/dileepkb/acs2015_census_tract_data_part2.csv
Table censusdata truncated.
Data loaded successfully. Elapsed Time: 0.3384 seconds
```

J. [MUST] Results

Use this table to present your results. List only the methods that you actually completed. List only load times for the full amount of Census data (not the small test sample of OR data created in part C). For each loading method, load both parts of the Census data and measure/report the load time for each part separately in the corresponding column of the table.

Method	Time to load part1	Time to load part2
D. Simple inserts	37.48 seconds	38.57 seconds
E. Drop Indexes and Constraints	34.83 seconds	35.5 seconds
F. Disable Autocommit	6.202 seconds	5.054 seconds
G. Use UNLOGGED table	7.97 seconds	7.95 seconds
H. Temp Table with memory tuning		
I. copy_from	0.3292 seconds	0.3384 seconds

K. [SHOULD] Observations

Use this section to record any observations about the various methods/techniques that you used for bulk loading of the USA Census data. Did you learn anything about why various loading approaches produce varying performance results?

Simple Inserts:

It took around 37 seconds for part 1 and 38 seconds for part 2 to load.

Because it needs to handle each action independently, this method which inserts each row one at a time ends up being the slowest.

Drop Indexes and Constraints:

Part 1 took around 34 seconds to load, while Part 2 took approximately 35 seconds.

Indexes and restrictions can be removed to expedite loading because the database won't need to maintain them when entering data. But thereafter, you have to recreate them.

Disable Autocommit:

Part 1 loaded in 6 seconds and Part 2 loads in around 5 seconds.

When autocommit is off, rows are added as a single transaction, which speeds up the process and lowers the cost associated with several commits.

Use UNLOGGED table:

It took around nine seconds for parts one and two to load. Although loading data into an unlogged table is faster than writing to the database log, you run the risk of losing data in the event of a system crash.

copy_from:

Part 1 took less than a second (0.3292 seconds) to load, while Part 2 took about the same amount of time(0.3384)

Because it loads data from a file in one go via a direct database function, this is by far the fastest

approach. It greatly accelerates the procedure and removes a large portion of the typical overhead.

To put it briefly, the fastest approaches are usually those that minimize steps and simplify the process. The most effective way to load data is to use specific methods like `copy_from` and disable `autocommit`.